

# B+ Trees, Splay Trees, Code

CptS 223 - Fall 2017 - Aaron Crandall



# Today's Agenda

- Announcements
- Humor of the day
- HW1 and solutions
- B+ Trees
- Splay trees (wacky!)
- Code demos

# Announcements





- Wednesday will be Red-Black trees
- Friday is all midterm review:
  - Will go over chapters 1-4 + Red-Black trees
  - I promise that there will be no attendance taken, it's a totally optional session
- Yes, we won our football game on Saturday
  - It wasn't really a question
  - I do give props to Leach for keeping #3 in after the fumble/interception strategy


# Thing of the day: Sony has patented a contact lens that is blink powered and records video

“The future is already here — it's just not very evenly distributed.”

-- William Gibson



# Linux in space!



“Spaceborne” Linux Supercomputer Starts Running In Space, Achieves 1 Teraflop Speed

<https://fossbytes.com/spaceborne/>

HPE’s Spaceborne Computer was launched into the space using SpaceX Dragon Spacecraft. This beast was launched as a result of a partnership between Hewlett Packard Enterprise (HPE) and NASA to find out how high-performance computers perform in space. Now, this supercomputer is fully installed and operational in ISS.


# Going over HW1

# B+ Trees!

- An M-ary tree (instead of a bin-ary tree)
  - Allows M children per node
- Data is stored at leaf nodes
  - Instead of at all nodes
- Designed to store data in blocks on disk instead of in RAM
  - Allows handling of larger data sets by optimising for slower read/write speeds
- Used very often for indexing in SQL databases
  - Gets sequential data access  $O(N)$  if leaf nodes link to each other

# The Memory Hierarchy

- CPU at the top
- Slowest data at bottom
- Handling disk accesses is what B-trees do



Source: Rambus




# Properties of a B-tree\* (really, a B+ tree)

1. The data items are stored at leaves.
2. The nonleaf nodes store up to  $M - 1$  keys to guide the searching; key  $i$  represents the smallest key in subtree  $i + 1$ .
3. The root is either a leaf or has between two and  $M$  children.
4. All nonleaf nodes (root can be smaller) have between  $\lceil M/2 \rceil$  and  $M$  children.
5. All leaves are at the same depth and have between  $\lceil L/2 \rceil$  and  $L$  data items.

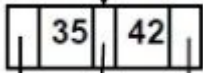
\* This is a B+ tree, an old skool B-tree stores data at all nodes, not just leaves

# Why the book talks about M & L

- There are two kinds of nodes in a B+ tree:
  - Index (or M) nodes of size M: The number of children a node can have
  - Data (or L) nodes of size L: The number of records a leaf can store
- M & L are calculated by how many you could stuff in a single file system block without going over.
  - Can only be whole numbers, since you can't store part of a node. Always round down.

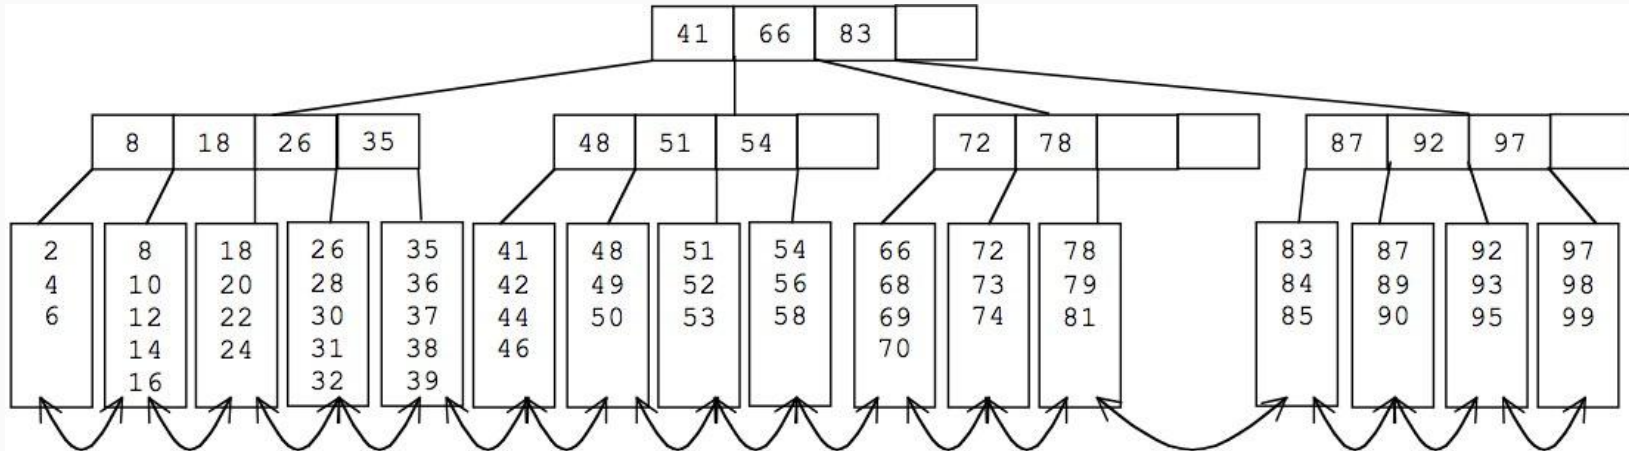


# M Nodes

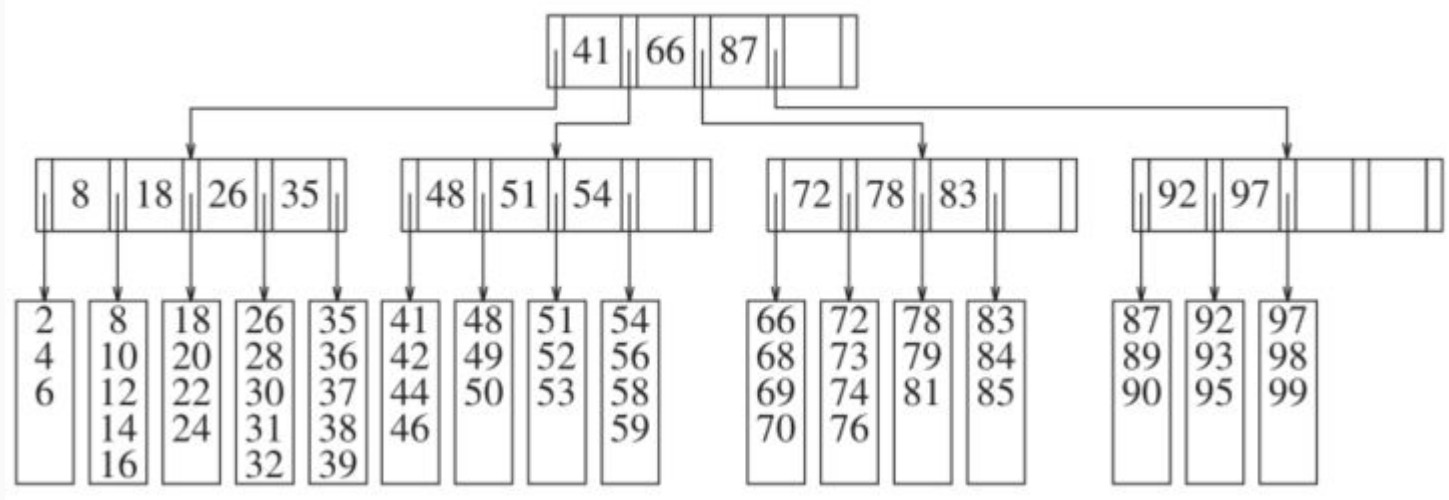
- Only store pointers to other nodes (M or L)
- Have M pointers and M-1 keys (how you index a record)
  - This example is a B-tree of order 3: 
- You calculate the order by how much physical space you use in a fs block:
  - $\text{fs blocksize} \geq M * \text{pointerbytes} + (M - 1) * \text{keysize}$
  - Round down for M
- Question: how big are pointers on your computer? (in bits)

# L nodes - only store data records (and maybe a pointer to the next L node)

- L nodes store the actual records
- The keys in M nodes are just the index value for records, not the data
- $fs \text{ blocksize} \geq L * \text{sizeof}(\text{record}) + \text{optional pointer}$

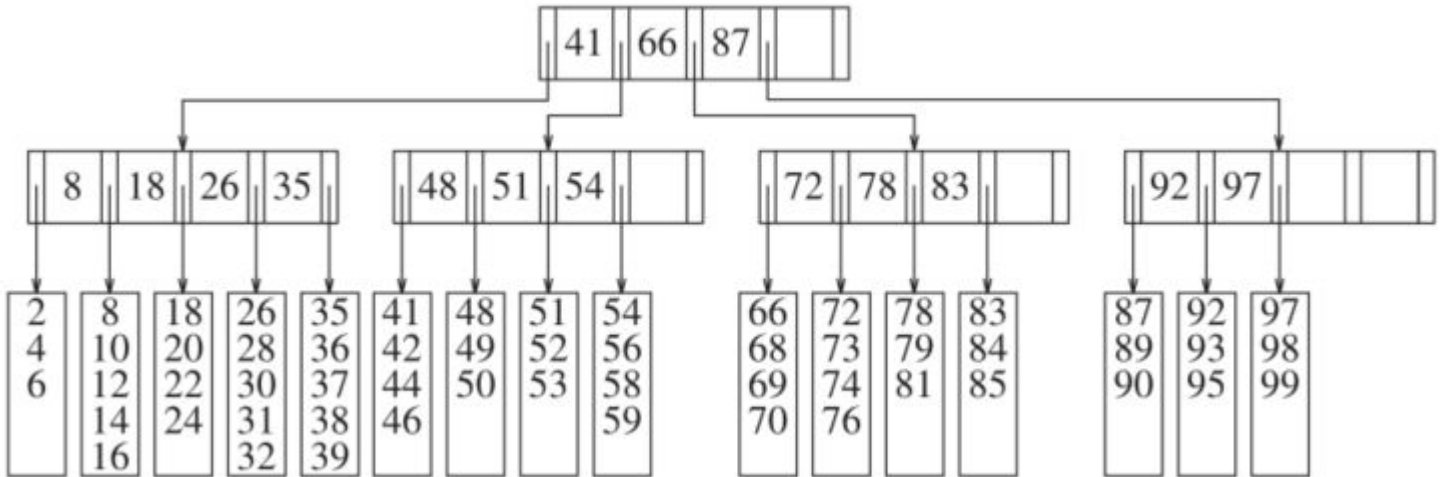


# B-Tree Example of order 5



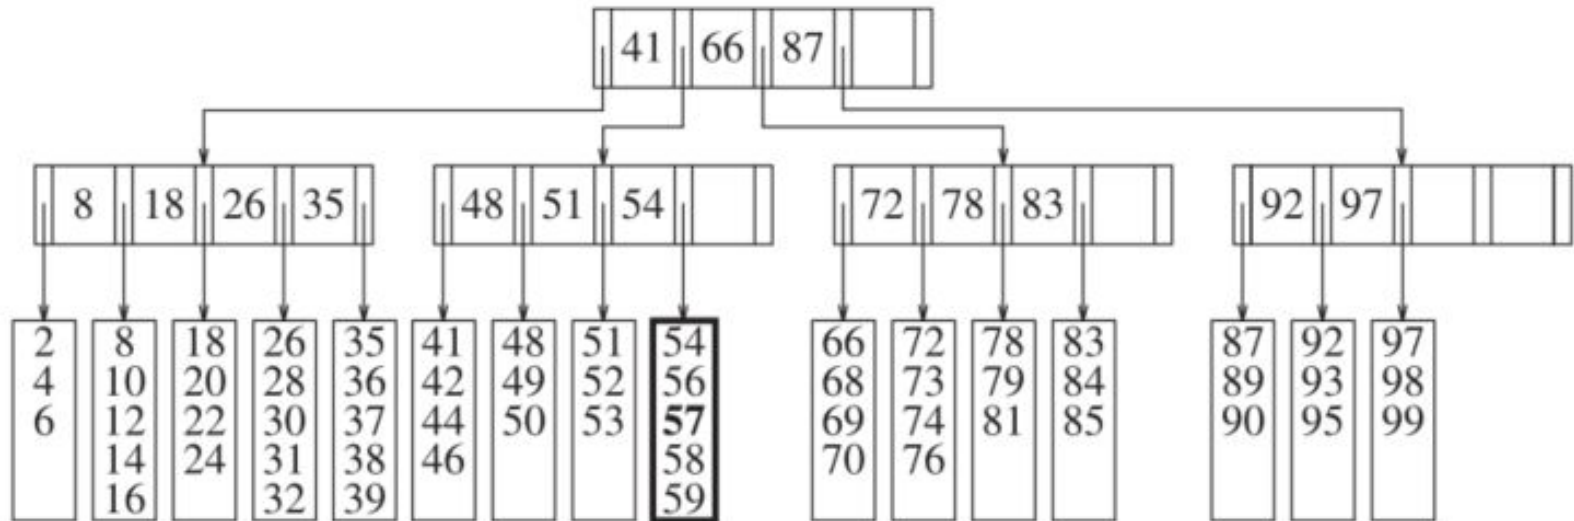
# B-tree: Insert, Exists, Delete

- Start at root, find pointer to next node (or leaf) based on keys, recurse
- Insert into leaf node if there's room: insert(57)



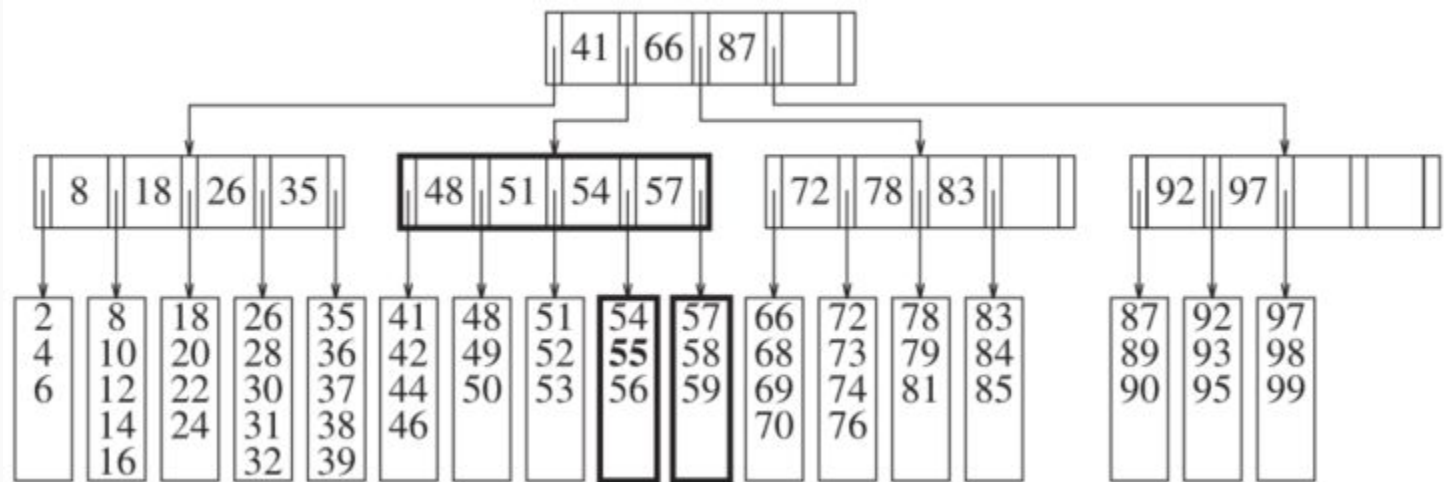
# insert(57) result

- There was room in the leaf node, so it just gets added
  - Have to read 3 blocks and write 1



# insert(55) causes a leaf split

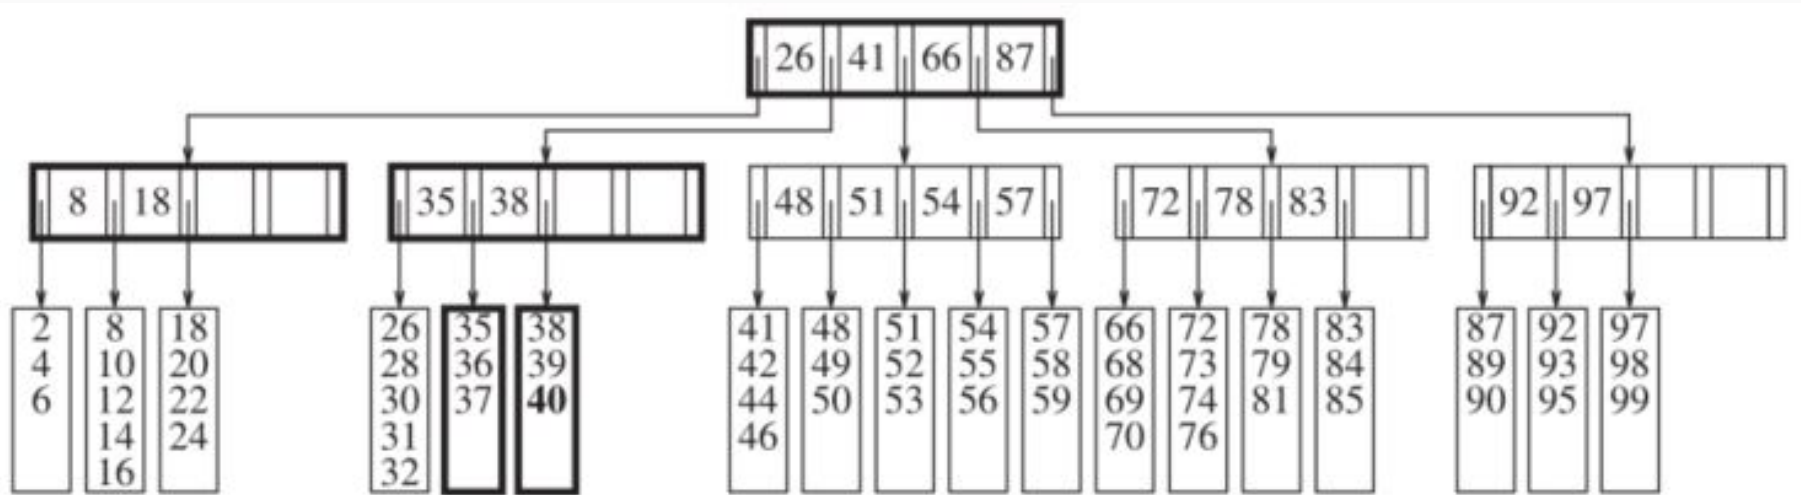
- insert(55) doesn't have room, so we split the leaf and update the node
- New leaves are guaranteed to have  $\lceil L/2 \rceil$  or greater elements (3 rd, 3 wr)





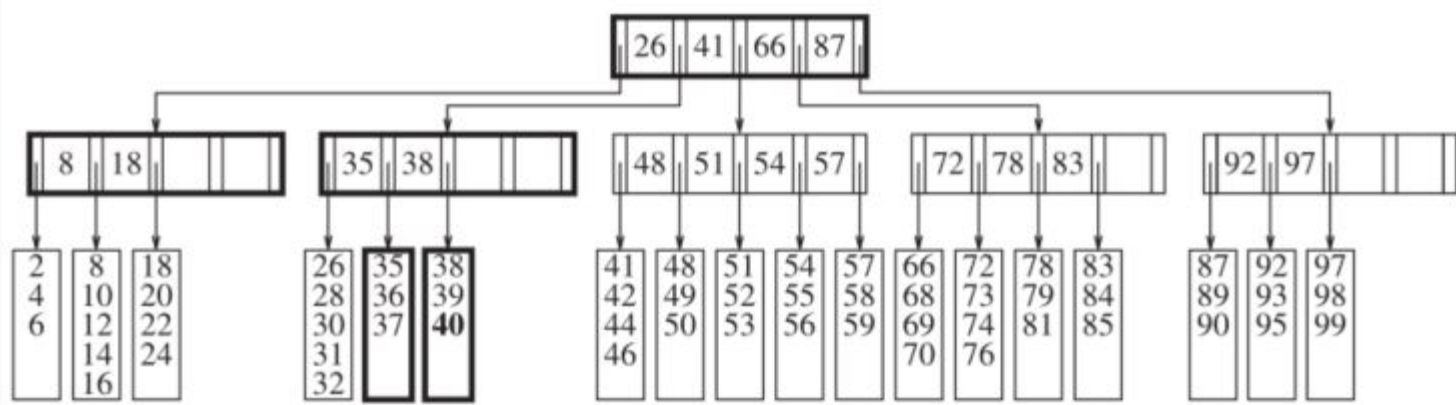
# insert(40) causes a node split

- insert(40) fills a leaf, then the parent node
- Parent splits recursively, up to root - 3 reads, then 5 writes

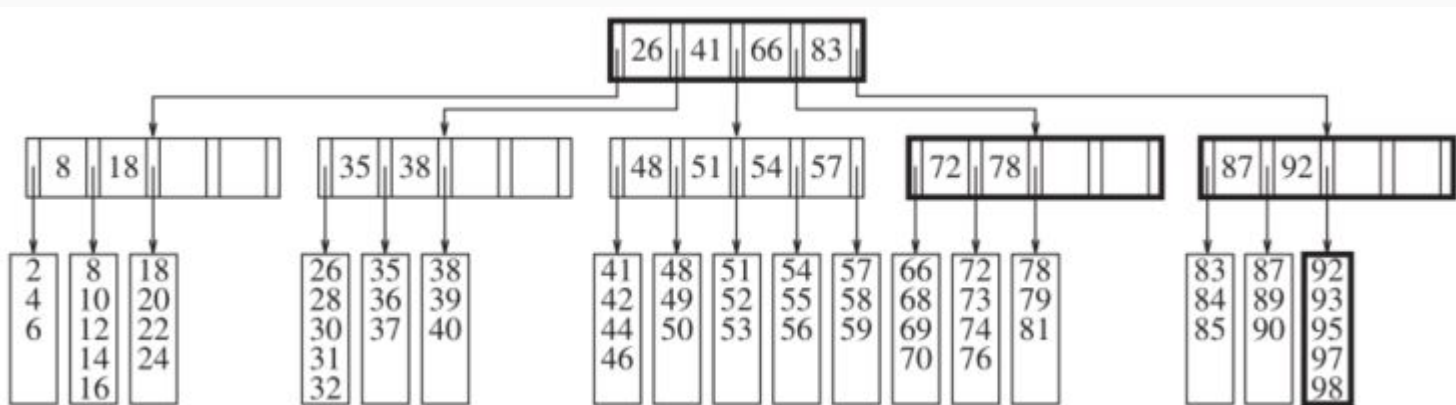


# delete

- Delete from leaf first, but...
- If leaf falls below minimum:
  - Borrow from node's other leaves (if they're not at the minimum)
  - Could merge with other node (if they are at the minimum)  $[M/2 + M/2]$  should have room
    - Can cause recursive node mergers, eventually shrinking the tree if root vanishes



Delete(99) causes a leaf to merge, then the parent to borrow a leaf from next door



# B-Tree summary

- Designed to align with filesystem nuances to speed up real-world searches
- Take more bookkeeping than other trees
- Allow for high speed in-order accesses if leaves are linked
- Take some more planning to fit with your hardware and file system
- Exploit the memory hierarchy whenever possible
- Used heavily in database implementations
- Visualization page:
  - <http://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>

# Sets & Maps in STL

- Sets: Can be used as vector or list, but has basic efficient searching
  - Internally, a binary tree
  - Takes hints to guess where to insert or search (normally local to last data accessed)
    - Which can give it  $O(1)$  access times if you're right. :-)
  - Normally implemented with top-down red black trees (see Chapter 12.2)
- Maps: Actually hashes internally
  - Uses a <key, value> for storing
  - Lookups done on key alone to retrieve value
  - See chapter 5, or lectures after the midterm
  - Key can be things like strings, so you could do lookups on names

# Splay Trees - An odd thing.

CptS 223 - Fall 2017 - Aaron Crandall



# Splay trees - Very different approach

- Book does this tree in passing. I had to suss out details here and there.
- “Simple” data structure, but only by less decision making
- Does guarantee no bad input sequences (akin to the BST linked list issue)
- Ends up with an amortized running time of  $O(M f(N)) \sim\sim O(\log(N))$  time
- **Every** time a node is accessed (created, read, modified), it is pulled to the root
- Exploits the locality of reference: Once you access a piece of data, you’re highly likely to access it again, or one right nearby
  - [https://en.wikipedia.org/wiki/Locality\\_of\\_reference](https://en.wikipedia.org/wiki/Locality_of_reference)

# Wait! *\*EVERY\** time a node is accessed?

- Yes, every time you read a node, you rotate and/or bring it to the root
- This means the trees is continually changing shape in significant ways
- It means we don't keep height or balance information because we'll always bring accessed nodes to the root, regardless of tree shape



# How to move the node to the root?


- Not by single rotations, it just makes the tree unbalanced in the other direction
- The book points out that a bad series of inserts:  $[1..N]$ , single rotations looks good, but gives  $\Omega(N^2)$  access times for in-order reads!
- So, instead of single rotations, we splay (zig-zag and zig-zig) instead
  - These kinds of rotations are used in Red-Black trees too

# Splaying

- Zig-Zag - The same as an AVL tree double rotation
- Zig-Zig - Rotating over the accessed node until it's the root


# Zig-Zag - Same as an AVL double rotate

- Look at grandparent of accessed node, find double rotation case, rotate!




# Zig-Zig - Pull accessed node to root over parent and grandparent

- Move X up to root over Parent and Grandparent



Then repeat until X is root of tree



# Prior example

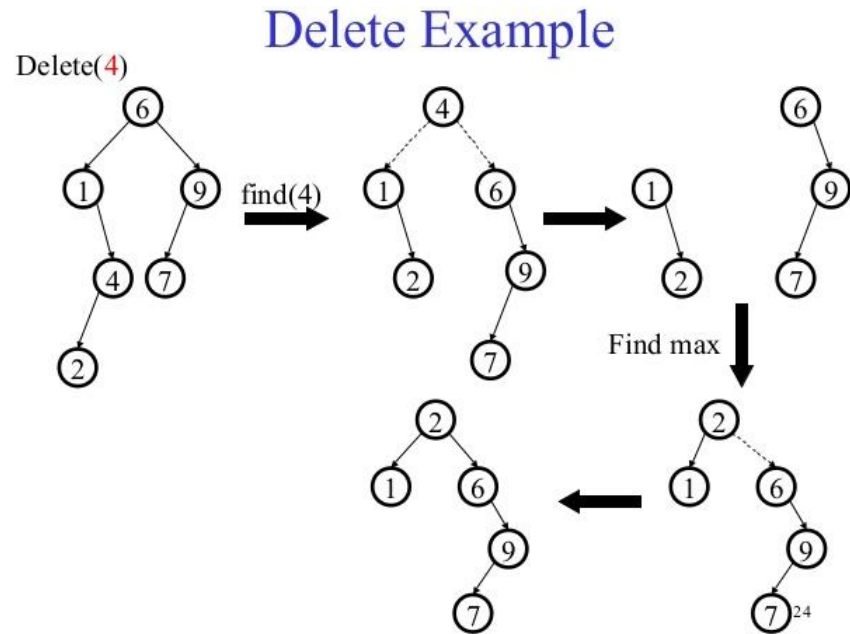
- First access of “1” takes  $O(N)$ , but then access of 2 takes  $\sim N/2$
- Nodes along the access path end up with their height about halved
- As this continues, the overall height ends up about  $\log(N)$

Result: When accesses are long, future accesses end up cheap

Concept of locality of reference: Most programs access data in a pattern based around either space or time. Splay trees help this by bringing relatively local values higher up the tree for the next access

# Deleting from a Splay Tree?

- 1) Access node to be deleted
  - a) Which brings it to the root
- 2) Treat children as two subtrees
- 3) Access max node in left tree
  - a) Which brings it to the root of the left tree, and it has no right child
- 4) Make the right tree the child of the new root of the left tree
- 5) Forget the node to be deleted
- 6) Done!



Lookin' at code - if there's time.



# AVL tree reminder

- Enforces balance of height difference of sub trees of no more than  $\pm 1$
- Access time of  $\log(N)$  by rotations on insert/delete
- L-L and R-R inserts take single rotations
- L-R and R-L inserts take double rotations

