

Sorting #2 - Shell and Heap

The Two Sorted Towers

CptS 223 - Fall 2017 - Aaron Crandall



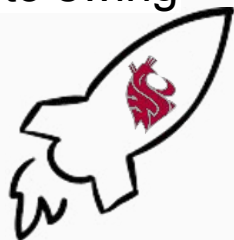
Today's Agenda

- Announcements
- Thing of the day
- Sorting is fun! - Doing a more formal heapsort
 - fginnorSstu!

Announcements



- Feedback from the Google hiring team that visited:
 - They were also quite vehement about the need for code reviews, version control, git, and exposure to UML. Students who don't understand graph traversal, recursion, hashes, and heaps are no-gos. Apparently they see a lot of students who go through these things!
 - So... notice that this is essentially our course (except the engineering aspects)?
- Hardware Hackathon is Saturday. Even if you don't participate, try to swing through and see what people are doing.



BONUS TOD:

Detecting Meter Maids with a RPi

- Guy parks on SF 2 hour limited street
- Camera & RPi to photo cars
- Motion detect for photo of cars & stuff
- Uses Tensor flow lib to image recog
- If 75%+ likely it's a meter maid:
 - Send SMS to phone via Twilio
 - Start clock on getting a ticket
- Bonus long term parking via tech
 - <http://peoplesparking.space/>



Why N^2 lower bound for simple sorts?

- First, define inversions:
an array of numbers is any ordered pair (i, j) having the property that:
 $i < j$ but $a[i] > a[j]$.
- Sort the example from last section: $\{34, 8, 64, 51, 32, 21\}$
 - I: $(34, 8), (34, 32), (34, 21), (64, 51), (64, 32), (64, 21), (51, 32), (51, 21)$, and $(32, 21)$
 - Exactly the number of swaps needed by insertion sort to sort the list
 - This is **always** the case since each swap fixes one inversion
 - A sorted list has no inversions, so it just needs checking: $O(N)$

So... given I inversions in N elements...

- There is always $O(N)$ work involved to test for an ordered list
- Then you add in I inversions, you get $O(I + N)$ for insertion sort
- Thus, you get $O(N)$ running time if there are no inversions
- How many inversions can you expect (on average)?

Calculating the average number of inversions in a list

- Assuming no duplicates
 - Otherwise, we need a way to estimate duplicates to subtract from total swaps
- Input becomes a permutation of the first N integers, all equally likely
 - Reverse our input list to get: 21, 32, 51, 64, 8, 34 = L_r
 - Consider every pair (x,y) with $y > x$
 - In exactly one of L and L_r , the ordered pair is an inversion
 - The total number of pairs in a list L and the reverse becomes: $N(N-1)/2$
 - The average is half of the total, so $N(N-1)/4$ inversions
- So... $O(N) + O([N^2-N]/4) \rightarrow \Omega(N^2)$ min total swaps are required to sort

Simple sorting algorithm minimum bound

Theorem:

Any algorithm that sorts by exchanging adjacent elements requires (N^2) time on average.

Proof:

The average number of inversions is initially $N(N-1)/4 = (N^2)/2$. Each swap removes only one inversion, so $(N^2)/2$ swaps are required.

Why is this cool?

- Example of a lower bound proof (it's a minimum work needed proof)
- Valid not only for insertion sort, but the entire class of sorting algorithms
 - Ones that only swap adjacent values
 - Also proves all future algorithms as yet undiscovered
- This proof is relatively simple, but in general lower bounds proofs are tough to do
- Shows that in order for a sorting algorithm to run subquadratic $o(N^2)$, it must do comparisons and exchanges between far apart elements.
 - Means it must eliminate more than one inversion per exchange

Shell sort (we might do it later in detail...)

- Named after Donald Shell
- Relatively simple algorithm using groups of numbers as sorting blocks
 - First to break the $O(N^2)$ barrier, with a $o(N^2)$ algorithm
 - Does far away comparisons, then shrinks down to insertion sort
 - Highly sensitive to your choices for how to shrink down the distance of the comparisons.
 - Proofs are complex, especially with the compound behavior
- Works great for 10's of thousands of elements
- Easy to code iteratively

Heapsort!

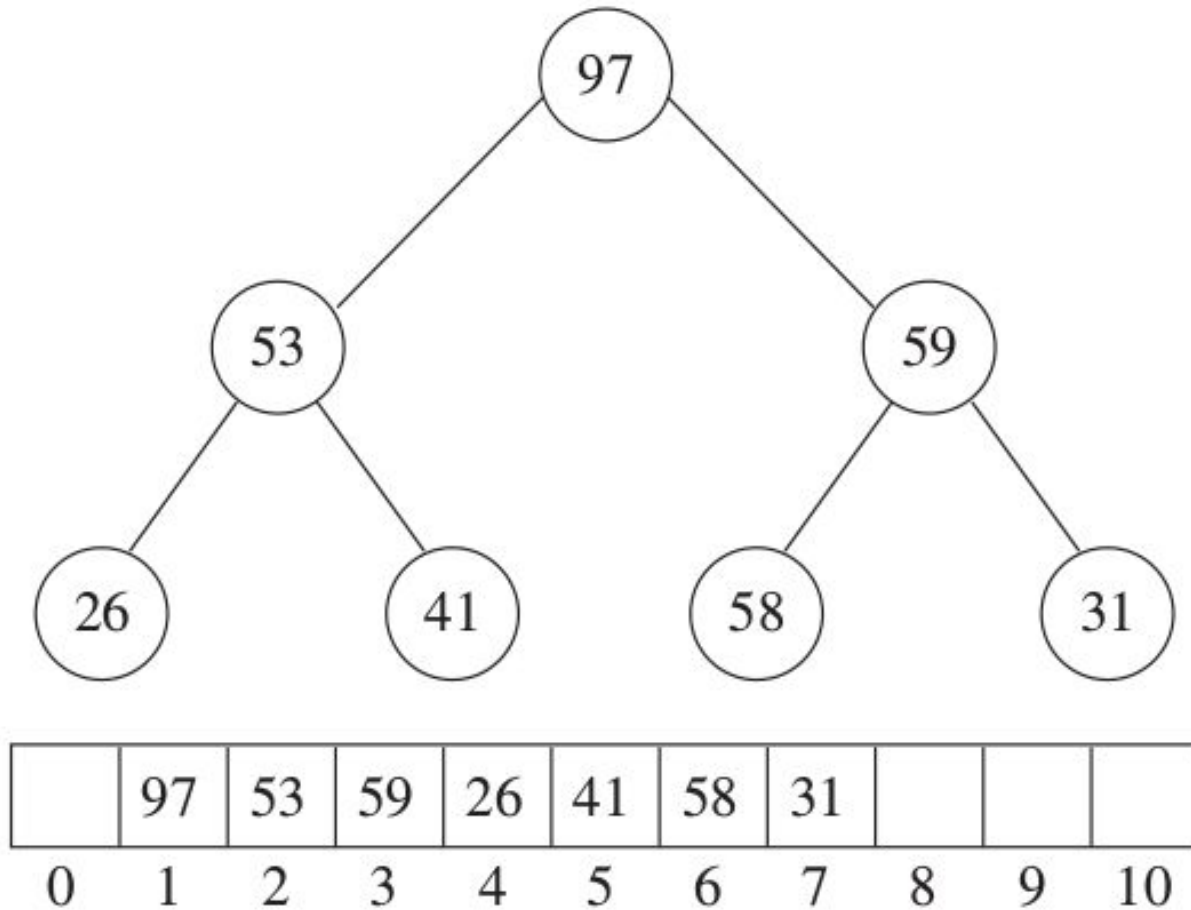
- Basically put your array in a heap, then pull it out at the root
- Works in $O(N \log N)$ time
- BuildHeap works in $O(N)$ time
- deleteMin() works in $O(\log N)$ time
 - We do this N times
- Total is: BuildHeap + $N * \text{deleteMin}$ $\rightarrow O(N) + O(N * \log(N))$
 - $O(N \log N)$
- Why not just roll with this and be done seeking more algorithms?

Because we sometimes forget memory complexity of an algorithm!

- Why is it a problem here?
- You need to store the sorted data in a second array (in the naive implementation), which doubles your memory needs. Space of: $O(2N)$
- Can be fixed by storing the sorted data in the same array as the heap shrinks down during deleteMin
 - Result is a reverse sorted array
 - Use a minHeap for $x > y$ sorting or a maxHeap for $y < x$ sorting

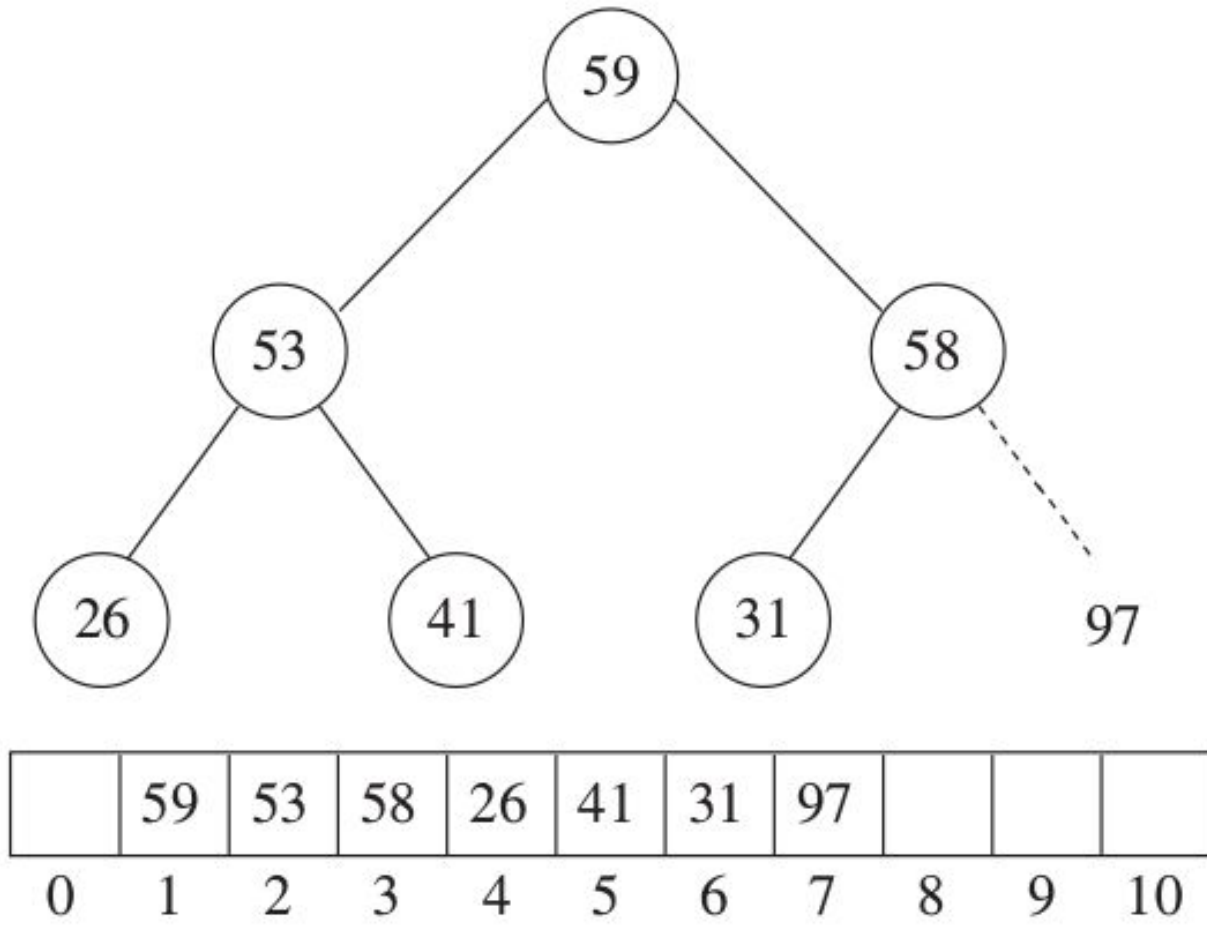
Example heap

- MaxHeap
- Took $O(N)$ to heapify



First iteration

- Took $O(\log N)$



Look at some code!

Note: it's starting at element 0, not 1.

Their visual example doesn't do this? Good one.

Monday: Mergesort

- Mergesort and Quicksort are the two most often used sorts in “real” implementations.
 - Mergesort is the standard in Java
 - Quicksort is the standard in C++

