

Sorting #5 - Final quicksort + Counting + Radix

CptS 223 - Fall 2017 - Aaron Crandall



Today's Agenda

- Announcements
- Thing of the day

Announcements



- Next MA MUST go out after class
- I'm so very done with this week, it was a long one!



Thing of the day: Waterfall Printer



Quicksort - Named (mostly) appropriately

- Runs in $O(N \log N)$ time
- Uses a divide and conquer strategy (like merge sort)
- Very sensitive to implementation
 - Pick a good pivot or you're doomed
- Recursive algorithm
 - Can work in the single array, so no doubling of memory space

Performance of Quicksort

- Average running time: $O(N \log N)$
- Worst case: $O(N^2)$
 - Can be made exponentially unlikely with a small tweak
 - All of the performance issues are centered around picking the pivot
- Often combined with other sorts to improve efficiency:
 - When N (of any sublist) is $5 \leq N \leq 20 \rightarrow$ change to Insertion sort


Basic Idea: Divide over pivot, repeat

- For list S:
 - If `S.size()` is 0 or 1, return
 - Pick element `v` in `S`, name it pivot
 - Partition `S - {v}` (take out the pivot) into two groups:
 - Those smaller than `v` -> `S_1`
 - Those bigger than `v` -> `S_2`
 - Return `{quicksort(S_1), v, quicksort(S_2)}`


WIN

Divide over pivot

- Pick pivot
- Divide into two lists
 - S₁ -> Those < pivot
 - S₂ -> Those > pivot



Combine back together



```
template <typename Comparable>
void SORT( vector<Comparable> & items )
{
    if( items.size( ) > 1 )
    {
        vector<Comparable> smaller;
        vector<Comparable> same;
        vector<Comparable> larger;

        auto chosenItem = items[ items.size( ) / 2 ];
```

```
for( auto & i : items )
{
    if( i < chosenItem )
        smaller.push_back( std::move( i ) );
    else if( chosenItem < i )
        larger.push_back( std::move( i ) );
    else
        same.push_back( std::move( i ) );
}
```

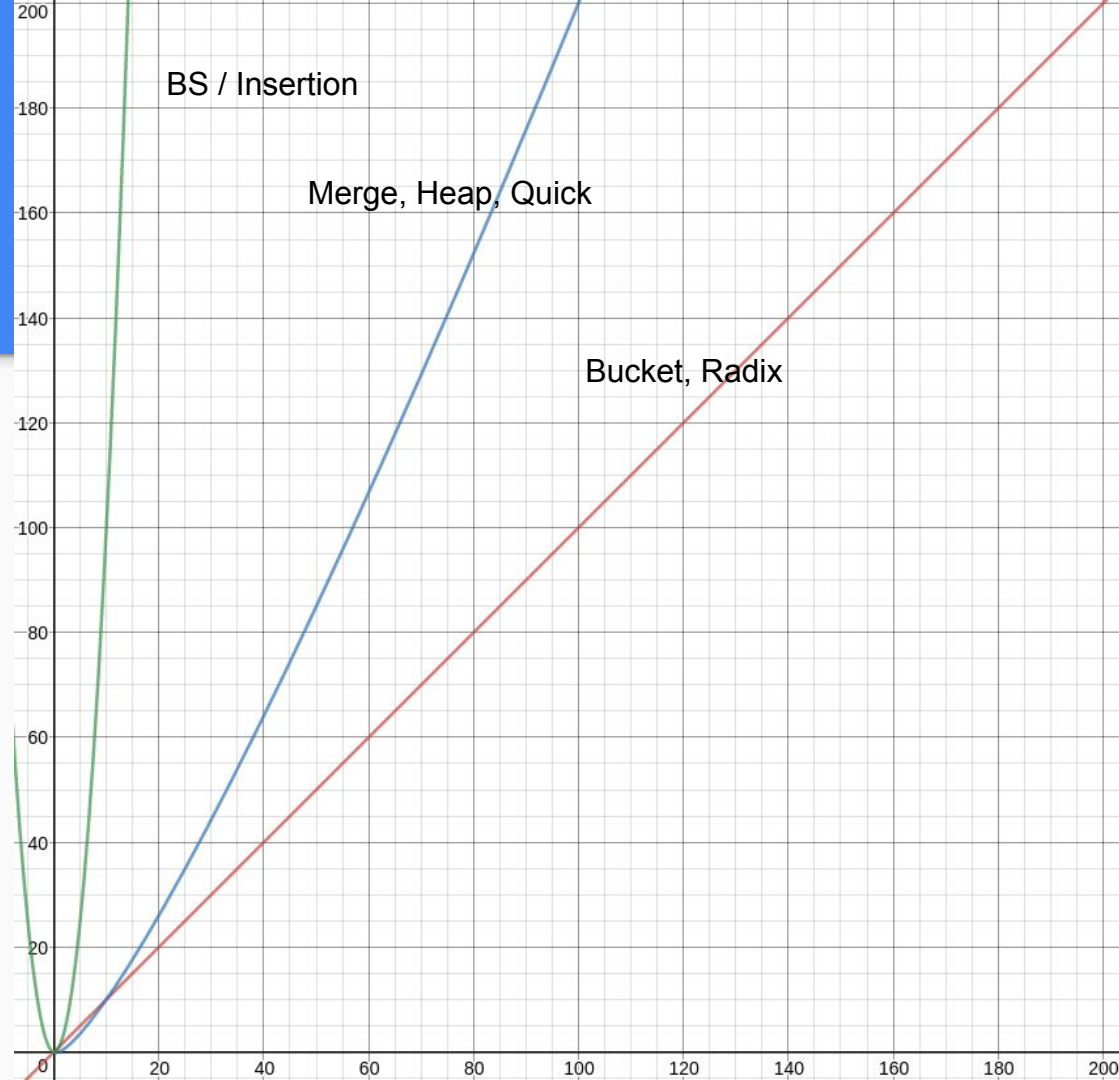
```
SORT( smaller );    // Recursive call!
SORT( larger );    // Recursive call!
```

```
std::move( begin( smaller ), end( smaller ), begin( items ) );
std::move( begin( same ), end( same ), begin( items ) + smaller.size( ) );
std::move( begin( larger ), end( larger ), end( items ) - larger.size( ) );
}

}
```


Linear Time Sorting

- Counting / Bucket sort
- Radix sort



Linear Time Sorts

- If your data has certain features, there's linear time sorting algorithms
 - $O(N) \ll O(N \log N)$
- Normally, this is exploiting size or ordering of data
- These are not “general purpose” sorting algorithms
 - Use some kind of non-comparison approach
- Ones in book:
 - Bucket (Counting) sort
 - Radix sort



Bucket (Counting) sort

- Linear time sort for small integers
- Given: $Arr = [A_1, A_2, \dots, A_N]$ of positive integers smaller than M .
- Make int counts[M], initialized to all 0.
- Read input: for each value in Arr you increment counts with: $counts[A_i]++$
- Then print out all non-zero buckets counts[i] times
- Results are found in $O(M+N)$ time
 - If $M == O(N)$, then final result is $O(N)$
- Power comes through M-way comparisons in $O(1)$ time
 - Similar to hashing algorithm

How useful is bucket sort?

- Requires having only positive integers
- Requires a known maximum value
- Requires an array of integers[maximum value] of storage space:
 $O(\text{maxValue in array to sort})$
- Does this perfect storm of features happen often?

YES

Consider some use cases

- Ages of people
- Social security number sorting: XXX-XX-XXXX -> counts[1000000000]
 - Too big? Unsure
- Car mileages in a fleet
- Could combine with a hash table, hashed by sorted key to lookup data?
 - Best of both worlds! Boom!
- Calculating medians of data sets

Different approach: Radix sort

- Historically used to sort punch cards, so also called card sort
- Can work on integers or strings
- Sorts by one indexed position at a time
 - Done right to left (won't work otherwise!) - must go from LSB to MSB
- Is a stable sort
- Effectively is a multi-pass bucket sort

Radix sort description

- Data must be N strings/values in range $(0..b^p-1)$
 - N == set of members in array
 - b == number of buckets, which is determined by the size of the alphabet in the data
 - p == number of passes, which is the max length of the strings/numbers
- Do passes over each digit or string position
 - Go from least significant to most significant
 - Put item into bucket as indexed by position (needs b buckets)
 - When done, read all items out in FIFO (to be a stable sort) order back to original array
 - Repeat p times, moving index from least to most significant position
- Result is done in $O(p(N+b))$ time!

Counting radix sort - sorting by numbers

INITIAL ITEMS:	064, 008, 216, 512, 027, 729, 000, 001, 343, 125
SORTED BY 1's digit:	000, 001, 512, 343, 064, 125, 216, 027, 008, 729
SORTED BY 10's digit:	000, 001, 008, 512, 216, 125, 027, 729, 343, 064
SORTED BY 100's digit:	000, 001, 008, 027, 064, 125, 216, 343, 512, 729

Each pass orders the items in a more significant way.

The results kind of pop into existence at the end.

Book code

- Not complex
- Requires the right *kind* of data set to be useful

```
void radixSortA( vector<string> & arr, int stringLen )
{
    const int BUCKETS = 256;
    vector<vector<string>> buckets( BUCKETS );

    for( int pos = stringLen - 1; pos >= 0; --pos )
    {
        for( string & s : arr )
            buckets[ s[ pos ] ].push_back( std::move( s ) );

        int idx = 0;
        for( auto & thisBucket : buckets )
        {
            for( string & s : thisBucket )
                arr[ idx++ ] = std::move( s );

            thisBucket.clear( );
        }
    }
}
```

What limits and why isn't it everywhere?

- Why doesn't this get used for long strings?
 - The key is in how string comparisons are done in programming languages
 - Also, has to handle a set with varying length strings (consider zero padded)
- Why doesn't this get used for large alphabets
 - How many total Unicode characters are there?
 - Book uses the extended ASCII set of 256, but...
 - Unicode has 17 planes of 65,536 each totalling 1,114,112 characters
 - Only 10% are allocated so far, which gives us around 115,000 for b.

External sorting - working on disks/tapes

- Like Trees, there's a need for sorting outside of main memory
 - B+ Trees were designed to solve this on disk
 - Sorting on tapes directly has a whole suite of options
 - Still relevant given the size of tapes: Sony's are up to 185 TB per tape
 - IBM/Sony just did 330 TB/tape last month: <https://goo.gl/n5CNds>
 - Companies still do many backups to tape even today
- Primary issues include:
 - Only linear access to items on the tape with any speed (forward or backward)
 - Incredibly slow compared to RAM if done out of order
 - Can use multiple tapes in parallel if your hardware supports it
- We're not going to spend too much time on this (actually, we're done)

Sorting summary

- The ordering of data is a huge research field
- Small advances can save companies huge money
- Picking the right algorithm to use needs to understand:
 - Data to sort
 - Primary consideration is input size
 - Pre-sorted, or almost sorted situations
 - Programming environment behaviors
 - Opportunity to do linear time sorting
- Most situations call for: Insertion sort, shellsort, mergesort, or quicksort
 - Be aware of opportunities for linear time sorting via bucket or radix sorts

Monday: Probably not doing graphs yet

- We have one kind of free day next week:
Monday: ??? - linux tools, more git for group work, etc
Wednesday: Midterm review
Friday: Midterm #2 - Heaps, Hashing, and Sorting

If I've gone fast enough, we can watch
hacker video for a while

<https://www.youtube.com/watch?v=hqKafI7Amd8>