

Graph Algorithms #1 - Introduction & Theory

CptS 223 - Fall 2017 - Aaron Crandall




Today's Agenda

- Announcements
- Thing of the day
- Introducing graph algorithms

Announcements



- Days left:
 - 8 more real days
 - 3 days in dead week
 - Final
- Targeting return of finals on Friday
- Will go over exam on Wednesday for logistical reasons
- Yes, we're so far behind on grading it's nail biting for me
 - It's got to be caught up ASAP



Thing of the Day


South Korea's first giant robot
take its first steps:

<https://www.youtube.com/watch?v=7rgFtkMiXms>

Jeff Bezos →



Graphs and Graph Algorithms







Wrong kind of graph

- What is a graph?



A graph $G = (V, E)$ consists of a set of vertices, V , and a set of edges, E .

- Each edge is a pair (v, w) , where $v, w \in V$.
- Graph algorithm?
 - An algorithm designed to work with data kept in a graph structure




Where have we seen graphs already?


Linked lists



Trees (Directed Acyclic Graph / DAG)



Hash tables (if you zoom out)



Even stacks


What isn't a graph?
Well... not much
(if anything)

How do we calculate big-O for graphs?

- In trees, heaps, hashing, stacks
 - Number of updates or operations to complete
 - Push == add node to head of list
 - Insert (tree) == traversal to bottom of tree, then node create & update
- In sorting:
 - Number of swaps/moves and comparisons done
- For graphs, it's based around two things:
 - Edges traversed, nodes (vertices) inspected
 - $|E|$ == number of edges || $|V|$ == number of vertices



Directed graphs - digraphs

- Directed graphs are when the edges are directed.
- This means they have a front and a back, normally shown as an arrow
- Where have we seen these already?
- $(v,w) \in E$ -- Edge from v to w




Undirected Graphs

- Edges do not have a front and back, normally shown with a line (no arrow)
 - Edges can be traversed in either direction
 - Think of it being the difference between one way streets and normal streets
 - Both nodes are adjacent to each other
- This below is (v, w) , which is also (w, v)




Adjacency

- Vertex w is adjacent to v if and only if $(v, w) \in E$
 - This means you can traverse the edge from v to w
- When using undirected edges, (v, w) means (w, v) so w and v are both adjacent to each other





W is *adjacent* to V



W is *adjacent* to V
 V is *adjacent* to W

Weight or Cost of an edge

- Edges can carry a cost to traverse them
 - For example, two intersections are connected and the cost is how many meters long the connecting road is



Degree of a vertex

- The number of adjacent vertices to the vertex
- *Indegree* is the number of incoming edges
- *Outdegree* is the number of outgoing edges



Paths



- A path is a sequence of vertices
 - $w_1, w_2, w_3, \dots, w_N$ such that $(w_i, w_{i+1}) \in E$ for $1 \leq i < N$
- The length of the path is the number of edges on the path (not vertices!)
 - So the length = $\text{count}(V) - 1$
- The path can go from a vertex to itself
 - If that path has no edges, it has a length of zero. This is a special case
- The path can do (v, v) , which is a loop
 - Normally loops don't happen in most algorithm traversals, but can happen
- Simple paths are where all vertices are distinct (no repeated vertices)
 - Exception: First and last can be the same if it's a path *and* a loop.

Cycles


- Directed graph of at least length 1
 - such that $w_1 == w_N$ (you get back to the start w)
- It's a simple cycle if the path is **simple**
 - No repeated nodes in the path, except the start & end
- In undirected graphs the edges must be **distinct**
 - You can only use an edge once in a cycle
- It's a **Directed Acyclic Graph** if it has no cycles
 - A special case called a DAG
 - Where have we seen these?




Connected vs. Disconnected

- An undirected graph is connected if there's a path from every vertex to every other vertex
- A directed graph like that is **strongly connected**
- If it's all included, but not strong then it's **weakly connected**
 - This means you have dead ends in the paths you can make
 - The graphs → are weakly connected

Connected Graph





Connected



Disconnected

Complete graph

- When there's an edge between every pair of vertices
 - There's a path of length 1 between every pair of vertices

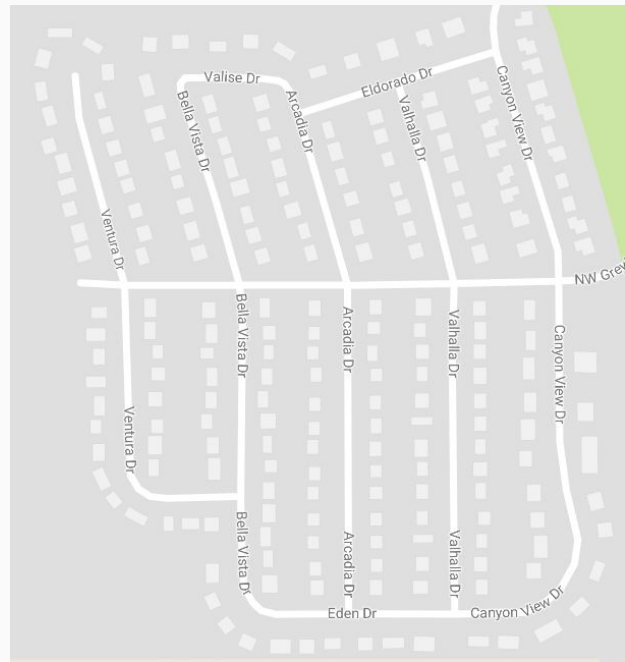



Examples of situations that suit graphs

- Airport connections
- Road trip route planning
- Traffic flow
- Networking
- LinkedIn
- Course prerequisites
- What else?
 - Bacon!


Storing and representing graphs

- Adjacency matrix
 - For each edge (u, v) , we set $A[u][v]$ to true ; otherwise the entry in the array is false
 - If edges have weights, set $A[u][v]$ equal to the weight and use either a very large or a very small weight as a sentinel to indicate nonexistent edges (INF, -INF)
 - Requires $\Theta(|V|^2)$ space, only appropriate if $|E| \sim \Theta(|V|^2)$
 - Example of badness with street maps






	a	b	c	d	e
a	0	1	1	1	0
b	1	0	0	1	0
c	1	0	0	0	0
d	1	1	0	0	1
e	0	0	0	1	0



	a	b	c	d	e
a	0	1	0	1	0
b	0	0	0	1	0
c	1	0	0	0	0
d	1	0	0	0	1
e	0	0	0	0	0








	0	1	2	3	4
0	0	2	3	0	0
1	2	0	15	2	0
2	3	15	0	0	13
3	0	2	0	0	9
4	0	0	13	9	0

**Adjacency Matrix Representation of
Weighted Graph**



Dense vs. Sparse

- Most graphs are ***sparse***
 - This means $|E| \ll |V|$
- Better solution for these is an adjacency list representation
 - Keep a list of all adjacent vertices for each vertex
 - Space requirement becomes $O(|E| + |V|)$
 - Instead of $\Theta(|V|^2)$ with the matrix
 - Weights can be kept with edges in adjacency list
 - Standard way to represent graphs






	a	b	c	d	e
a	0	1	1	1	0
b	1	0	0	1	0
c	1	0	0	0	0
d	1	1	0	0	1
e	0	0	0	1	0




	a	b	c	d	e
a	0	1	0	1	0
b	0	0	0	1	0
c	1	0	0	0	0
d	1	0	0	0	1
e	0	0	0	0	0



Direct comparison of adjacency matrix vs. adjacency list

If time, then do topo sort!


- Topological sort is an ordering of vertices in a directed acyclic graph, such that if there is a path from v_i to v_j , then v_j appears after v_i in the ordering
- Can't work if there's a cycle in the graph
- Does not guarantee a unique ordering
- Used for deciding scheduling of work units
 - Edges represent the dependency of work units
 - Only those with an indegree of 0 can be done next



Topographical Sorting Example

$\{v_1, v_2, v_5, v_4, v_3, v_7, v_6\}$ and $\{v_1, v_2, v_5, v_4, v_7, v_3, v_6\}$ are both valid topological orderings

- 1) Find node with no in edges
 - a) Indegree of zero
 - b) Called a "source node"
- 2) Print out node && remove from graph
- 3) Repeat



For Wednesday:

Reviewing exam

More on graphs, starting into Dijkstra's algorithm to stage for next MA